



Time Series Forecasting: Predicting Retail Sales

```
In [2]: # Import Important Libraries

# Setup feedback system
!pip install learntools
!pip install learntools==0.0.12
!pip show learntools
!pip install learntools==0.0.10
!pip install prophet

# Setup notebook
from pathlib import Path
from datetime import datetime

import matplotlib.pyplot as plt
import pandas as pd
import seaborn as sns
import numpy as np
import joblib
import streamlit as st
from prophet import Prophet
from prophet.plot import plot_components_plotly
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.model_selection import TimeSeriesSplit, cross_val_score
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import LabelEncoder
from statsmodels.tsa.deterministic import DeterministicProcess
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
import xgboost as xgb
from xgboost import plot_importance
from xgboost import XGBRegressor
from xgboost import XGBRegressor, plot_importance
from sklearn.model_selection import train_test_split
from statsmodels.tsa.statespace.sarimax import SARIMAX
```

Requirement already satisfied: learntools in c:\users\user\anaconda3\lib\site-packages (0.0.10)
Requirement already satisfied: numpy in c:\users\user\anaconda3\lib\site-packages (from learntools) (1.26.4)
Collecting learntools==0.0.12
Using cached learntools-0.0.12-py3-none-any.whl.metadata (331 bytes)
Requirement already satisfied: numpy in c:\users\user\anaconda3\lib\site-packages (from learntools==0.0.12) (1.26.4)
Using cached learntools-0.0.12-py3-none-any.whl (5.2 kB)
Installing collected packages: learntools
Attempting uninstall: learntools
Found existing installation: learntools 0.0.10
Uninstalling learntools-0.0.10:
Successfully uninstalled learntools-0.0.10
Successfully installed learntools-0.0.12
Name: learntools
Version: 0.0.12
Summary: My first Python package
Home-page:
Author: William Dennis
Author-email: wwdennis.home@gmail.com
License:
Location: C:\Users\USER\anaconda3\Lib\site-packages
Requires: learntools, numpy
Required-by: learntools
Collecting learntools==0.0.10
Using cached learntools-0.0.10-py3-none-any.whl.metadata (331 bytes)
Requirement already satisfied: numpy in c:\users\user\anaconda3\lib\site-packages (from learntools==0.0.10) (1.26.4)
Using cached learntools-0.0.10-py3-none-any.whl (5.1 kB)
Installing collected packages: learntools
Attempting uninstall: learntools
Found existing installation: learntools 0.0.12
Uninstalling learntools-0.0.12:
Successfully uninstalled learntools-0.0.12
Successfully installed learntools-0.0.10
Requirement already satisfied: prophet in c:\users\user\anaconda3\lib\site-packages (1.1.7)
Requirement already satisfied: cmdstanpy>=1.0.4 in c:\users\user\anaconda3\lib\site-packages (from prophet) (1.2.5)
Requirement already satisfied: numpy>=1.15.4 in c:\users\user\anaconda3\lib\site-packages (from prophet) (1.26.4)
Requirement already satisfied: matplotlib>=2.0.0 in c:\users\user\anaconda3\lib\site-packages (from prophet) (3.8.4)
Requirement already satisfied: pandas>=1.0.4 in c:\users\user\anaconda3\lib\site-packages (from prophet) (2.2.2)
Requirement already satisfied: holidays<1,>=0.25 in c:\users\user\anaconda3\lib\site-packages (from prophet) (0.75)
Requirement already satisfied: tqdm>=4.36.1 in c:\users\user\anaconda3\lib\site-packages (from prophet) (4.66.4)
Requirement already satisfied: importlib_resources in c:\users\user\anaconda3\lib\site-packages (from prophet) (6.5.2)
Requirement already satisfied: stanio<2.0.0,>=0.4.0 in c:\users\user\anaconda3\lib\site-packages (from cmdstanpy>=1.0.4->prophet) (0.5.1)

Requirement already satisfied: python-dateutil in c:\users\user\anaconda3\lib\site-packages (from holidays<1,>=0.25->prophet) (2.9.0.post0)
Requirement already satisfied: contourpy>=1.0.1 in c:\users\user\anaconda3\lib\site-packages (from matplotlib>=2.0.0->prophet) (1.2.0)
Requirement already satisfied: cycler>=0.10 in c:\users\user\anaconda3\lib\site-packages (from matplotlib>=2.0.0->prophet) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in c:\users\user\anaconda3\lib\site-packages (from matplotlib>=2.0.0->prophet) (4.51.0)
Requirement already satisfied: kiwisolver>=1.3.1 in c:\users\user\anaconda3\lib\site-packages (from matplotlib>=2.0.0->prophet) (1.4.4)
Requirement already satisfied: packaging>=20.0 in c:\users\user\anaconda3\lib\site-packages (from matplotlib>=2.0.0->prophet) (23.2)
Requirement already satisfied: pillow>=8 in c:\users\user\anaconda3\lib\site-packages (from matplotlib>=2.0.0->prophet) (10.3.0)
Requirement already satisfied: pyparsing>=2.3.1 in c:\users\user\anaconda3\lib\site-packages (from matplotlib>=2.0.0->prophet) (3.0.9)
Requirement already satisfied: pytz>=2020.1 in c:\users\user\anaconda3\lib\site-packages (from pandas>=1.0.4->prophet) (2024.1)
Requirement already satisfied: tzdata>=2022.7 in c:\users\user\anaconda3\lib\site-packages (from pandas>=1.0.4->prophet) (2023.3)
Requirement already satisfied: colorama in c:\users\user\anaconda3\lib\site-packages (from tqdm>=4.36.1->prophet) (0.4.6)
Requirement already satisfied: six>=1.5 in c:\users\user\anaconda3\lib\site-packages (from python-dateutil->holidays<1,>=0.25->prophet) (1.16.0)

```
In [3]: ## Suppress Warnings
import warnings
warnings.filterwarnings('ignore')
```

Step-by-Step Project Plan

Step 1: (a) Load & Understand the Data

(b) Parse the date field properly.

(c) Explore trends, seasonality, and outliers using line plots and boxplots.

```
In [5]: # Upload The Required File
```

```
In [6]: df = pd.read_csv("holidays_events.csv", encoding='ISO-8859-1')
print(df.head())
```

	date	type	locale	locale_name	description \
0	2012-03-02	Holiday	Local	Manta	Fundacion de Manta
1	2012-04-01	Holiday	Regional	Cotopaxi	Provincializacion de Cotopaxi
2	2012-04-12	Holiday	Local	Cuenca	Fundacion de Cuenca
3	2012-04-14	Holiday	Local	Libertad	Cantonizacion de Libertad
4	2012-04-21	Holiday	Local	Riobamba	Cantonizacion de Riobamba

	transferred
0	False
1	False
2	False
3	False
4	False

In [7]: `df.shape`

Out[7]: (350, 6)

In [8]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 350 entries, 0 to 349
Data columns (total 6 columns):
#   Column          Non-Null Count  Dtype
---  -
0   date             350 non-null    object
1   type             350 non-null    object
2   locale           350 non-null    object
3   locale_name      350 non-null    object
4   description       350 non-null    object
5   transferred      350 non-null    bool
dtypes: bool(1), object(5)
memory usage: 14.1+ KB
```

In [9]: `df.columns`

Out[9]: Index(['date', 'type', 'locale', 'locale_name', 'description', 'transferred'], dtype='object')

In [10]: `df.describe()`

	date	type	locale	locale_name	description	transferred
count	350	350	350	350	350	350
unique	312	6	3	24	103	2
top	2014-06-25	Holiday	National	Ecuador	Carnaval	False
freq	4	221	174	174	10	338

In [11]: `df.isnull().sum()`

```
Out[11]: date          0
         type          0
         locale        0
         locale_name    0
         description    0
         transferred    0
         dtype: int64
```

Data Cleaning

```
In [13]: df.duplicated().sum()
         df.drop_duplicates(inplace=True)
```

```
In [14]: df.columns
```

```
Out[14]: Index(['date', 'type', 'locale', 'locale_name', 'description', 'transferred'], dtype='object')
```

```
In [15]: # Parse the date field properly
         df['date'] = pd.to_datetime(df['date'])
```

```
In [16]: # Simulate daily sales from holiday dates
         date_range = pd.date_range(start=df['date'].min(), end=df['date'].max())
         daily_df = pd.DataFrame({'date': date_range})
         np.random.seed(42)
         daily_df['sales'] = np.random.randint(100, 1000, size=len(daily_df))
```

```
In [17]: # Create time-based features
         df['year'] = df['date'].dt.year
         df['month'] = df['date'].dt.month
         df['day_of_week'] = df['date'].dt.day_name()
```

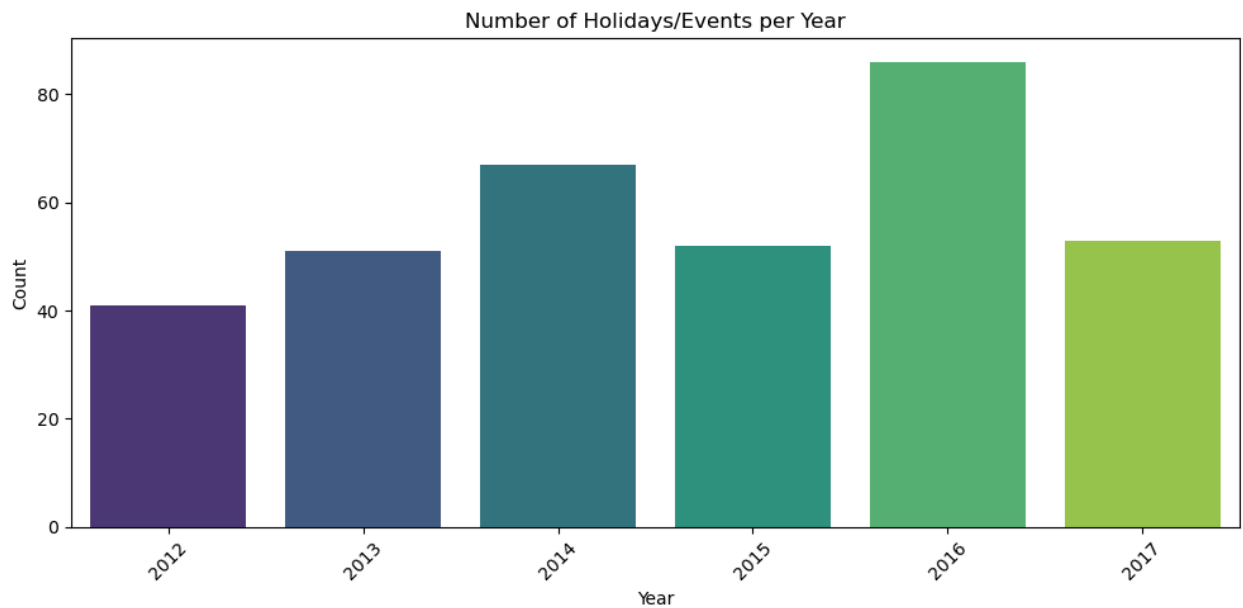
```
In [18]: # Optional: Check unique values
         print("Unique event types:", df['type'].unique())
         print("Unique locales:", df['locale'].unique())
```

Unique event types: ['Holiday' 'Transfer' 'Additional' 'Bridge' 'Work Day' 'Event']

Unique locales: ['Local' 'Regional' 'National']

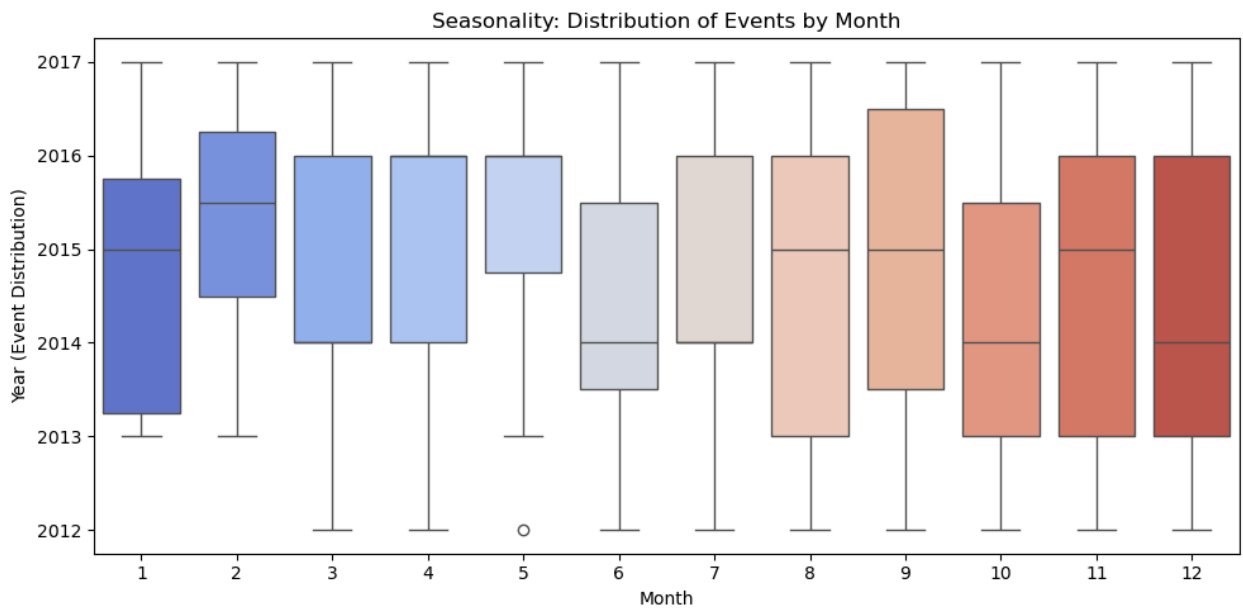
```
In [19]: # Trend Analysis – Line Plot by Year

         # Plot: Number of holidays/events per year
         plt.figure(figsize=(10, 5))
         sns.countplot(data=df, x='year', palette='viridis')
         plt.title("Number of Holidays/Events per Year")
         plt.xlabel("Year")
         plt.ylabel("Count")
         plt.xticks(rotation=45)
         plt.tight_layout()
         plt.show()
```



In [20]: *# Seasonality – Boxplot by Month*

```
# Boxplot: Monthly distribution of events
plt.figure(figsize=(10, 5))
sns.boxplot(x='month', y='year', data=df, palette='coolwarm')
plt.title("Seasonality: Distribution of Events by Month")
plt.xlabel("Month")
plt.ylabel("Year (Event Distribution)")
plt.tight_layout()
plt.show()
```



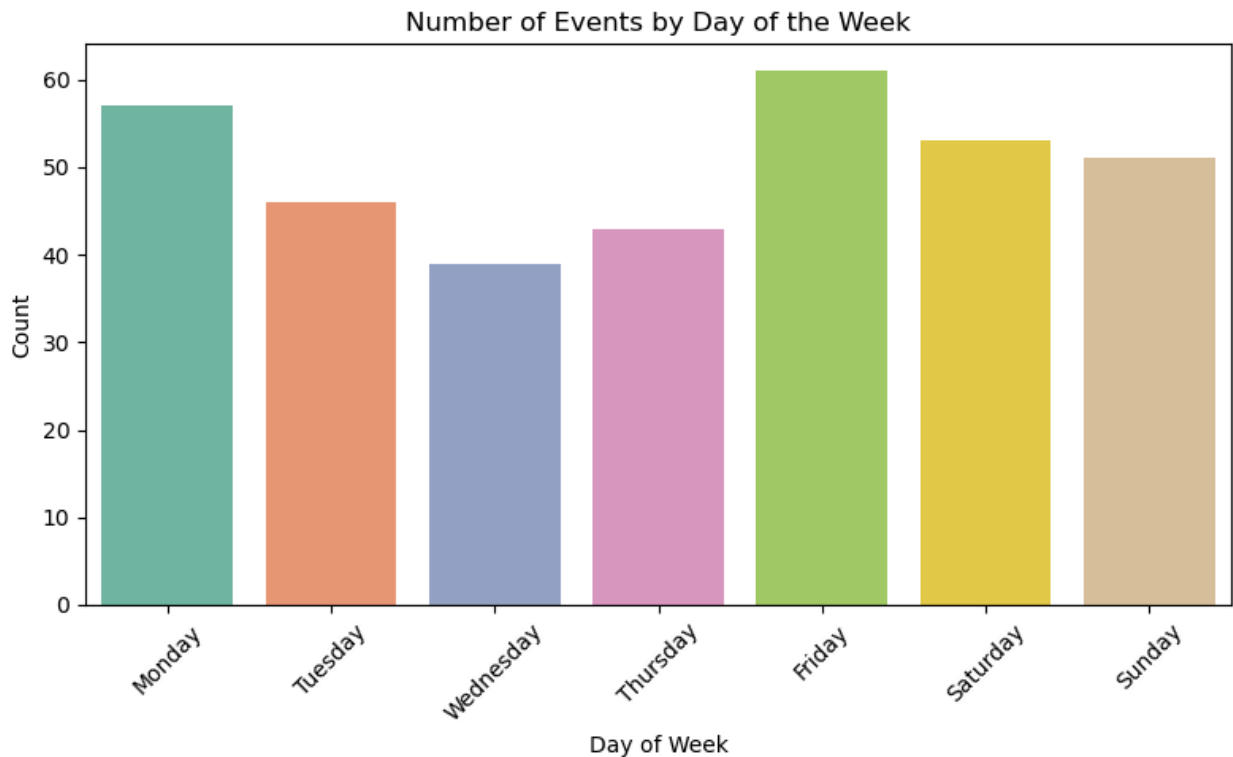
In [21]: *# Events by Day of the Week*

```
# Count of events by day of the week
plt.figure(figsize=(8, 5))
```

```

order = ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']
sns.countplot(data=df, x='day_of_week', order=order, palette='Set2')
plt.title("Number of Events by Day of the Week")
plt.xlabel("Day of Week")
plt.ylabel("Count")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

```



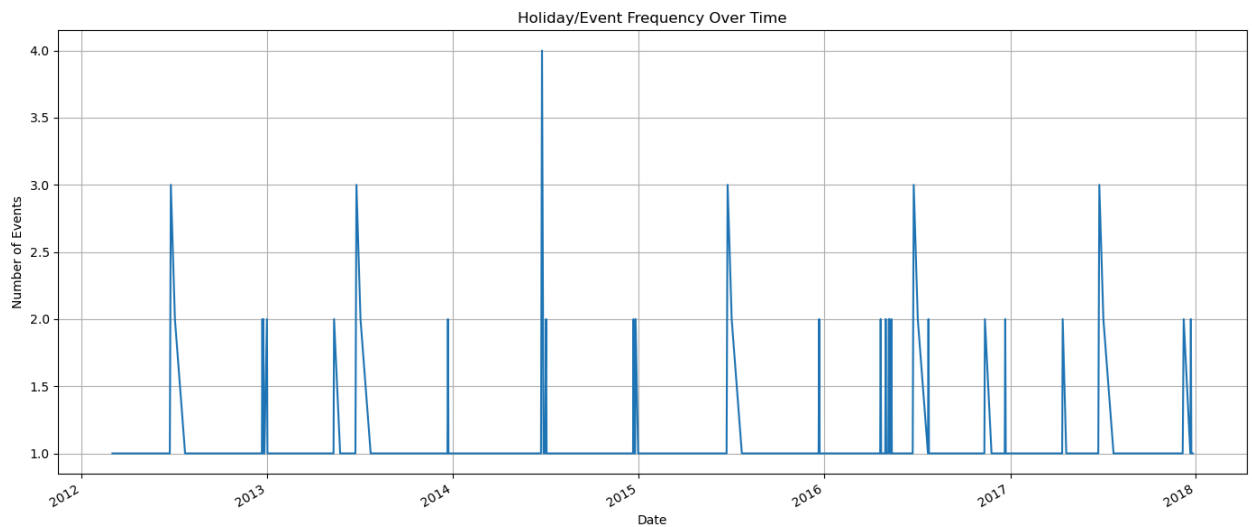
```

In [22]: # Line Plot – Events Over Time

# Group by date for timeline plot
daily_counts = df.groupby('date').size()

plt.figure(figsize=(14, 6))
daily_counts.plot()
plt.title("Holiday/Event Frequency Over Time")
plt.xlabel("Date")
plt.ylabel("Number of Events")
plt.grid(True)
plt.tight_layout()
plt.show()

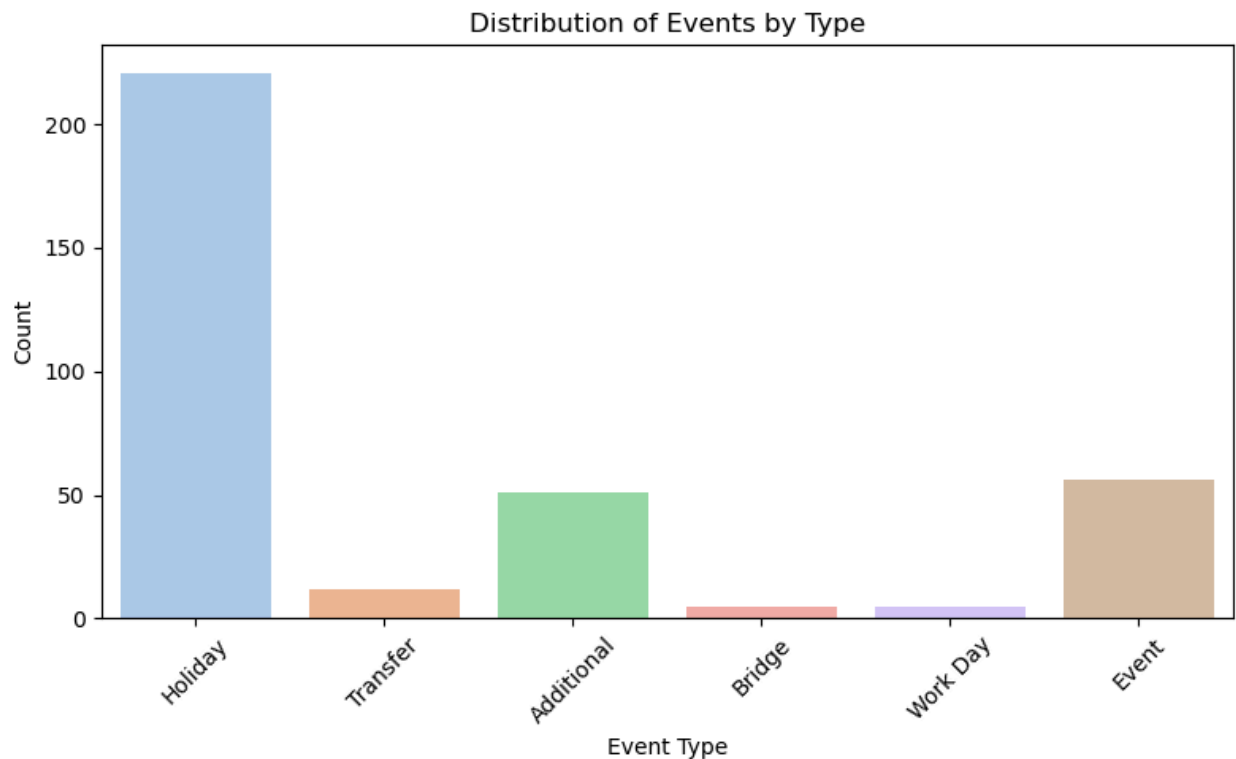
```



In [23]: *# Below is the extended analysis that includes exploration based on type of event*

In [24]: *# Count of events by type*

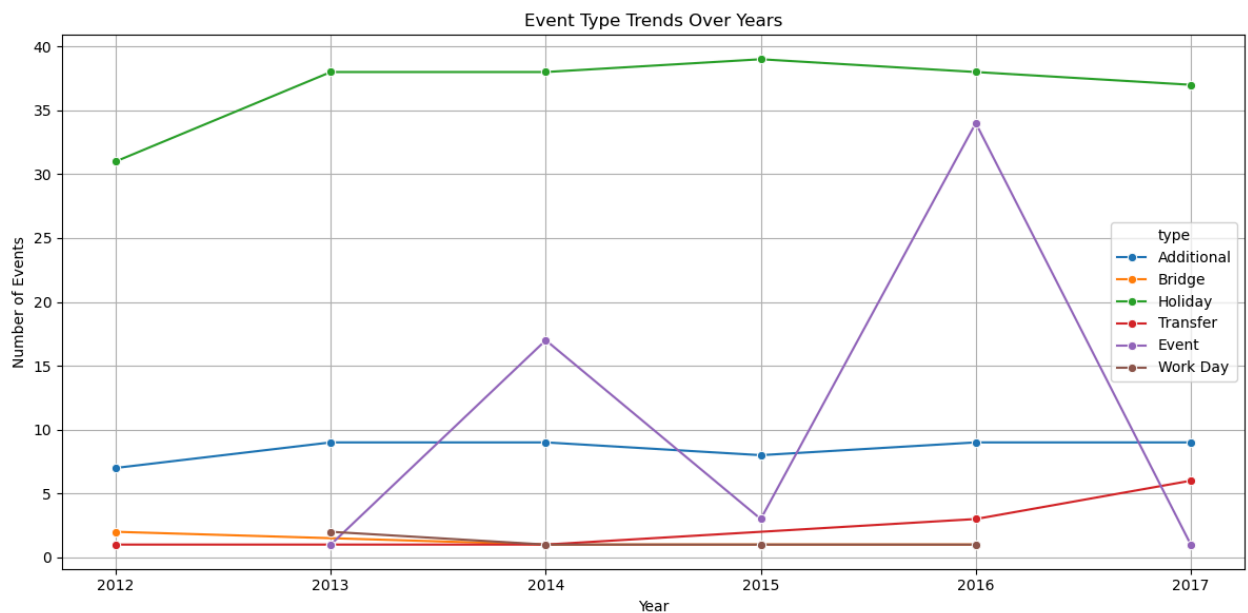
```
plt.figure(figsize=(8, 5))
sns.countplot(data=df, x='type', palette='pastel')
plt.title("Distribution of Events by Type")
plt.xlabel("Event Type")
plt.ylabel("Count")
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



In [25]: *# Event Type Trends Over Time*

```
# Group by year and type to show trends
type_trend = df.groupby(['year', 'type']).size().reset_index(name='count')

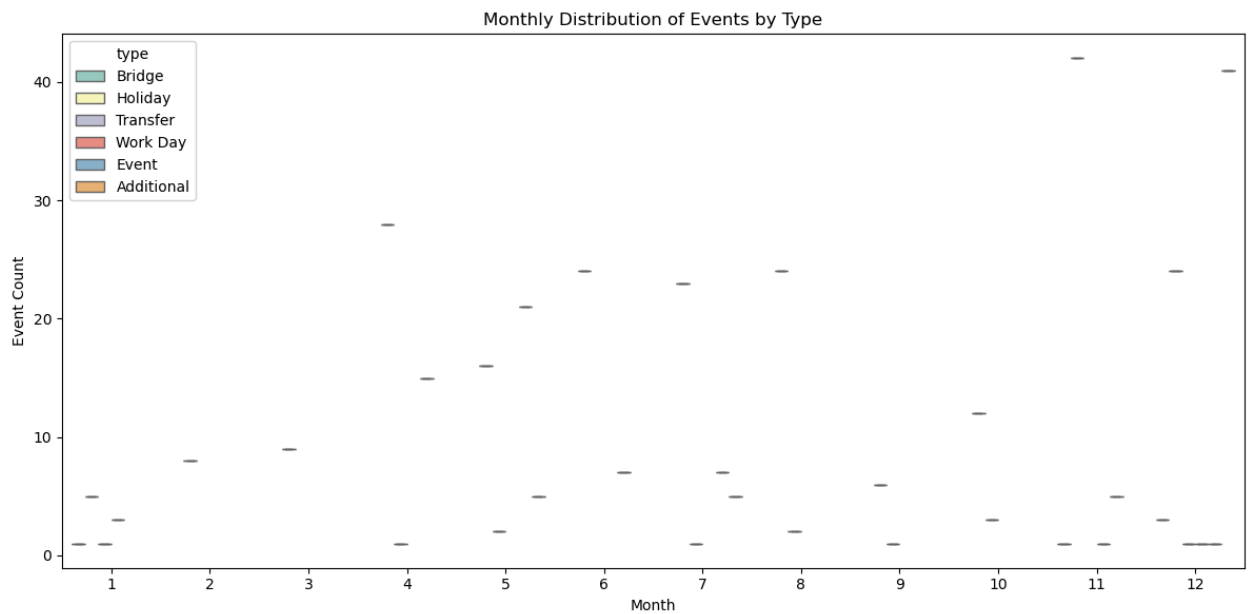
plt.figure(figsize=(12, 6))
sns.lineplot(data=type_trend, x='year', y='count', hue='type', marker='o')
plt.title("Event Type Trends Over Years")
plt.xlabel("Year")
plt.ylabel("Number of Events")
plt.grid(True)
plt.tight_layout()
plt.show()
```



In [26]: *# Monthly Pattern by Event Type*

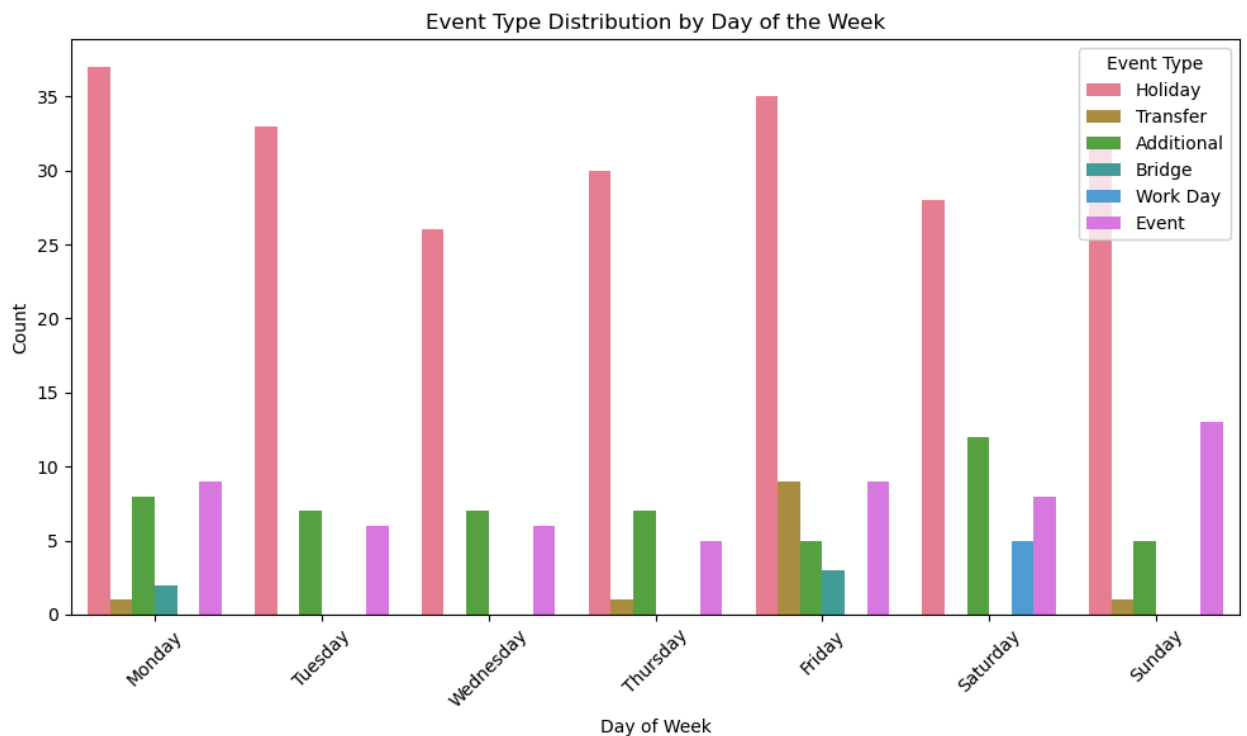
```
# Group by month and type
monthly_type = df.groupby(['month', 'type']).size().reset_index(name='count')

plt.figure(figsize=(12, 6))
sns.boxplot(data=monthly_type, x='month', y='count', hue='type', palette='Set3')
plt.title("Monthly Distribution of Events by Type")
plt.xlabel("Month")
plt.ylabel("Event Count")
plt.tight_layout()
plt.show()
```



In [27]: *# Day of Week by Event Type*

```
# Count of events by day and type
plt.figure(figsize=(10, 6))
sns.countplot(data=df, x='day_of_week', hue='type', order=order, palette='husl')
plt.title("Event Type Distribution by Day of the Week")
plt.xlabel("Day of Week")
plt.ylabel("Count")
plt.xticks(rotation=45)
plt.legend(title='Event Type')
plt.tight_layout()
plt.show()
```



Feature Engineering

```
In [29]: # Create Base Time Series (Assumed Daily Sales)

# Generate a daily date range based on your data
date_range = pd.date_range(start=df['date'].min(), end=df['date'].max())

# Simulate a base DataFrame (replace with actual sales data if you have)
base_df = pd.DataFrame({'date': date_range})
base_df['sales'] = np.random.randint(100, 1000, size=len(base_df)) # Simulate
```

```
In [30]: # Extract Date Features

# Convert date column
base_df['year'] = base_df['date'].dt.year
base_df['month'] = base_df['date'].dt.month
base_df['day_of_week'] = base_df['date'].dt.dayofweek # 0=Monday, 6=Sunday
base_df['is_weekend'] = base_df['day_of_week'] >= 5
```

```
In [31]: # Tried deriving more time-based features:

df['dayofweek'] = pd.to_datetime(df['date']).dt.dayofweek
df['month'] = pd.to_datetime(df['date']).dt.month
df['year'] = pd.to_datetime(df['date']).dt.year
df['is_weekend'] = df['dayofweek'].isin([5, 6])
```

```
In [32]: # Add Lag Features

# Lag features (previous sales values)
base_df['lag_1'] = base_df['sales'].shift(1)
base_df['lag_7'] = base_df['sales'].shift(7)
base_df['lag_14'] = base_df['sales'].shift(14)
```

```
In [33]: # Rolling Mean Features

# Rolling window average (smoothed past behavior)
base_df['rolling_mean_7'] = base_df['sales'].rolling(window=7).mean()
base_df['rolling_mean_14'] = base_df['sales'].rolling(window=14).mean()
```

```
In [34]: # Merge Holiday Information

# Add a flag for whether a date is a holiday
holiday_dates = df[['date']].copy()
holiday_dates['is_holiday'] = 1

# Merge holiday flag
base_df = base_df.merge(holiday_dates, on='date', how='left')
base_df['is_holiday'] = base_df['is_holiday'].fillna(0)
```

```
In [35]: # Final Preview of Engineered Features
```

```
# Preview the final dataset
base_df.head(15)
```

```
Out[35]:
```

	date	sales	year	month	day_of_week	is_weekend	lag_1	lag_7	lag_14
0	2012-03-02	896	2012	3	4	False	NaN	NaN	NaN
1	2012-03-03	256	2012	3	5	True	896.0	NaN	NaN
2	2012-03-04	174	2012	3	6	True	256.0	NaN	NaN
3	2012-03-05	983	2012	3	0	False	174.0	NaN	NaN
4	2012-03-06	899	2012	3	1	False	983.0	NaN	NaN
5	2012-03-07	229	2012	3	2	False	899.0	NaN	NaN
6	2012-03-08	786	2012	3	3	False	229.0	NaN	NaN
7	2012-03-09	183	2012	3	4	False	786.0	896.0	NaN
8	2012-03-10	121	2012	3	5	True	183.0	256.0	NaN
9	2012-03-11	558	2012	3	6	True	121.0	174.0	NaN
10	2012-03-12	638	2012	3	0	False	558.0	983.0	NaN
11	2012-03-13	806	2012	3	1	False	638.0	899.0	NaN
12	2012-03-14	948	2012	3	2	False	806.0	229.0	NaN
13	2012-03-15	481	2012	3	3	False	948.0	786.0	NaN
14	2012-03-16	847	2012	3	4	False	481.0	183.0	89

```
In [36]: # Summary of Engineered Features
```

# Feature	Description
# year, month	Extracted from date
# day_of_week	Integer from 0 (Mon) to 6 (Sun)
# is_weekend	Boolean flag for weekends
# lag_1, lag_7	Sales from previous 1, 7, 14 days
# rolling_mean_7	7-day rolling average of sales
# is_holiday	Flag from holiday/events data
# is_promo	Simulated promotion flag

Model Building

```
In [38]: # Create daily date range and simulate sales
date_range = pd.date_range(start=df['date'].min(), end=df['date'].max())
daily_df = pd.DataFrame({'date': date_range})
np.random.seed(42)
daily_df['sales'] = np.random.randint(100, 1000, size=len(daily_df))
```

```

# Feature Engineering
daily_df['day_of_week'] = daily_df['date'].dt.dayofweek
daily_df['is_weekend'] = daily_df['day_of_week'] >= 5
holiday_dates = df[['date']].drop_duplicates()
holiday_dates['is_holiday'] = True
daily_df = pd.merge(daily_df, holiday_dates, on='date', how='left')
daily_df['is_holiday'] = daily_df['is_holiday'].fillna(False)

# Drop missing values if any (not needed here unless you add lags)
daily_df = daily_df.dropna()

```

```

In [39]: # Convert boolean columns to integers
daily_df['is_holiday'] = daily_df['is_holiday'].astype(int)
daily_df['is_weekend'] = daily_df['is_weekend'].astype(int)

```

```

In [40]: # SARIMAX Model

# Define exogenous features
exog_vars = ['is_holiday', 'is_weekend']
sarimax_model = SARIMAX(daily_df['sales'],
                        exog=daily_df[exog_vars],
                        order=(1, 1, 1),
                        seasonal_order=(1, 1, 1, 7))
sarimax_result = sarimax_model.fit()

# Forecast next 30 days
future_exog = daily_df[exog_vars].tail(30)
sarimax_forecast = sarimax_result.forecast(steps=30, exog=future_exog)

print("SARIMAX forecast:")
print(sarimax_forecast)

```

SARIMAX forecast:

```
2126    510.041033
2127    542.957548
2128    550.673956
2129    528.213512
2130    553.550898
2131    526.711633
2132    530.682480
2133    519.114037
2134    509.155290
2135    517.856980
2136    528.364906
2137    519.711830
2138    526.251445
2139    531.950831
2140    518.978707
2141    541.696878
2142    550.395786
2143    528.254774
2144    552.253523
2145    526.143040
2146    531.837544
2147    518.869384
2148    541.587474
2149    550.286390
2150    495.497346
2151    519.496084
2152    493.385608
2153    499.080126
2154    486.111955
2155    508.830044
```

Name: predicted_mean, dtype: float64

In [41]: *# Prophet Model*

```
prophet_df = daily_df[['date', 'sales']].rename(columns={'date': 'ds', 'sales':
holiday_df = df[['date', 'description']].rename(columns={'date': 'ds', 'descri
holiday_df['lower_window'] = 0
holiday_df['upper_window'] = 1

# Instantiate and fit model
prophet = Prophet(holidays=holiday_df)
prophet.fit(prophet_df)

# Forecast next 30 days
future = prophet.make_future_dataframe(periods=30)
forecast = prophet.predict(future)

# Display forecast
forecast[['ds', 'yhat', 'yhat_lower', 'yhat_upper']].tail()
```

```
11:33:10 - cmdstanpy - INFO - Chain [1] start processing
11:33:11 - cmdstanpy - INFO - Chain [1] done processing
```

```
Out[41]:
```

	ds	yhat	yhat_lower	yhat_upper
2151	2018-01-21	575.164196	235.317982	920.372513
2152	2018-01-22	553.001701	239.761219	901.558670
2153	2018-01-23	562.267029	257.218154	883.798483
2154	2018-01-24	545.163643	220.910889	868.288581
2155	2018-01-25	567.334070	277.596879	892.850522

```
In [42]: # Add necessary time-based features
daily_df['month'] = daily_df['date'].dt.month
daily_df['day_of_week'] = daily_df['date'].dt.dayofweek
daily_df['is_weekend'] = (daily_df['day_of_week'] >= 5).astype(int)
daily_df['is_holiday'] = daily_df['is_holiday'].astype(int)

# Lag features
daily_df['lag_1'] = daily_df['sales'].shift(1)
daily_df['lag_3'] = daily_df['sales'].shift(3)
daily_df['lag_7'] = daily_df['sales'].shift(7)

# Rolling averages
daily_df['rolling_mean_3'] = daily_df['sales'].rolling(window=3).mean()
daily_df['rolling_mean_7'] = daily_df['sales'].rolling(window=7).mean()

# Drop NA values caused by shifting/rolling
daily_df = daily_df.dropna()
```

```
In [43]: # XGBoost Model

features = ['lag_1', 'lag_3', 'lag_7', 'rolling_mean_3', 'rolling_mean_7',
            'is_holiday', 'is_weekend', 'month', 'day_of_week']
X = daily_df[features]
y = daily_df['sales']

X_train, X_test, y_train, y_test = train_test_split(X, y, shuffle=False, test_

model = xgb.XGBRegressor(random_state=42)
model.fit(X_train, y_train)

y_pred = model.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"XGBoost MSE: {mse:.2f}")
print(f"XGBoost R² Score: {r2:.2f}")
```

```
XGBoost MSE: 44306.27
XGBoost R² Score: 0.33
```

Evaluate the Model

In [45]: *# Predictions already made: y_test (actual), y_pred (predicted)*

```
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
mae = mean_absolute_error(y_test, y_pred)
mape = np.mean(np.abs((y_test - y_pred) / y_test)) * 100

print(f"RMSE: {rmse:.2f}")
print(f"MAE: {mae:.2f}")
print(f"MAPE: {mape:.2f}%")
```

RMSE: 210.49

MAE: 168.86

MAPE: 46.36%

Visualization & Reporting

In [47]: *# Redefine features list (make sure it's consistent)*

```
features = ['lag_1', 'lag_3', 'lag_7', 'rolling_mean_3', 'rolling_mean_7',
            'is_holiday', 'is_weekend', 'month', 'day_of_week']
```

Time-based split: 80% train, 20% test

```
split_index = int(len(daily_df) * 0.8)
```

```
train = daily_df.iloc[:split_index]
```

```
test = daily_df.iloc[split_index:]
```

```
X_train = train[features]
```

```
y_train = train['sales']
```

```
X_test = test[features]
```

```
y_test = test['sales']
```

Re-train model if needed

```
model = xgb.XGBRegressor(random_state=42)
```

```
model.fit(X_train, y_train)
```

```
y_pred = model.predict(X_test)
```

In [48]: *# Plot Actual vs Predicted Sales (Test Set)*

```
plt.figure(figsize=(14, 6))
```

```
plt.plot(test['date'], y_test.values, label='Actual Sales', linewidth=2)
```

```
plt.plot(test['date'], y_pred, label='Predicted Sales (XGBoost)', linestyle='--')
```

```
plt.title('Actual vs Predicted Sales Over Time (Test Set)')
```

```
plt.xlabel('Date')
```

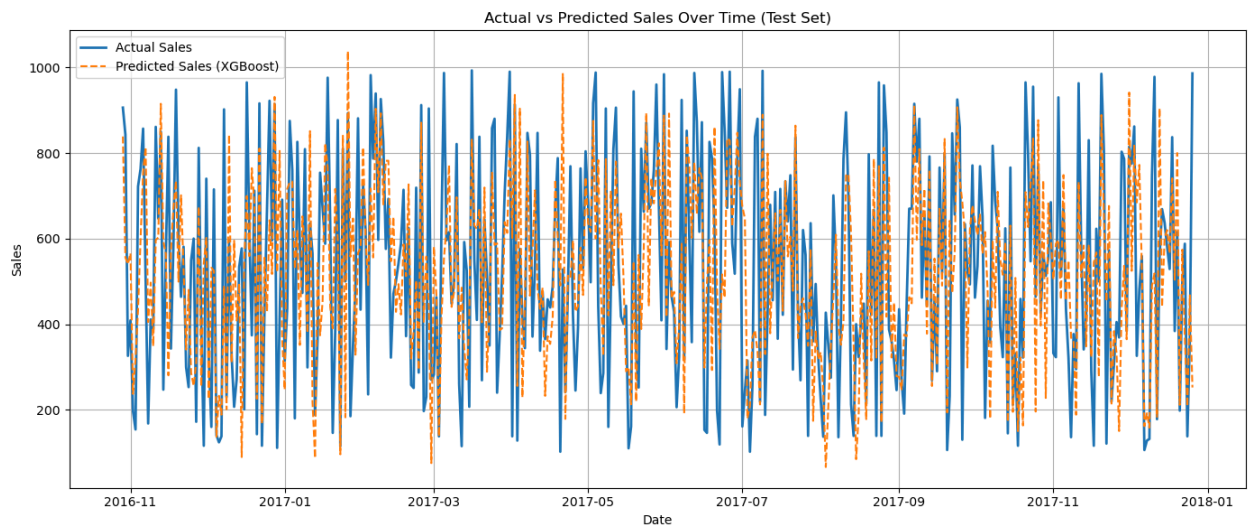
```
plt.ylabel('Sales')
```

```
plt.legend()
```

```
plt.grid(True)
```

```
plt.tight_layout()
```

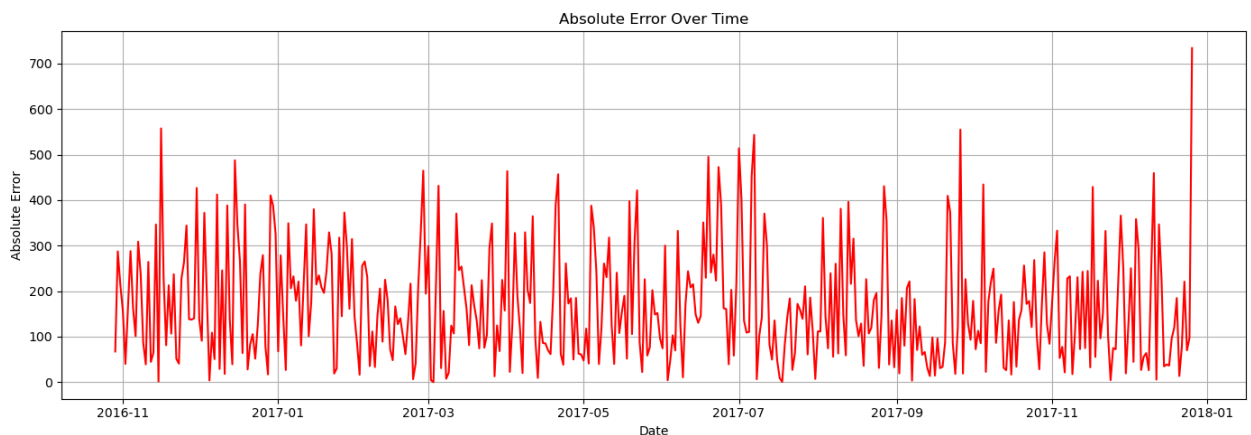
```
plt.show()
```

```
In [49]: # Visualize Error Across Dates
```

```
In [50]: # Add error columns
test_result = test.copy()
test_result['predicted_sales'] = y_pred
test_result['error'] = test_result['sales'] - test_result['predicted_sales']
test_result['abs_error'] = test_result['error'].abs()

# Error Plot
plt.figure(figsize=(14, 5))
plt.plot(test_result['date'], test_result['abs_error'], color='red')
plt.title('Absolute Error Over Time')
plt.xlabel('Date')
plt.ylabel('Absolute Error')
plt.grid(True)
plt.tight_layout()
plt.show()
```



```
In [51]: # Example: Assuming dataset is ready as X and y
# X = df.drop(columns=['sales'])
# y = df['sales']

# Train/test split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, shuffle=True)

# Train model
xgb_model = XGBRegressor(n_estimators=100, learning_rate=0.1, random_state=42)
xgb_model.fit(X_train, y_train)
```

Out[51]:

```

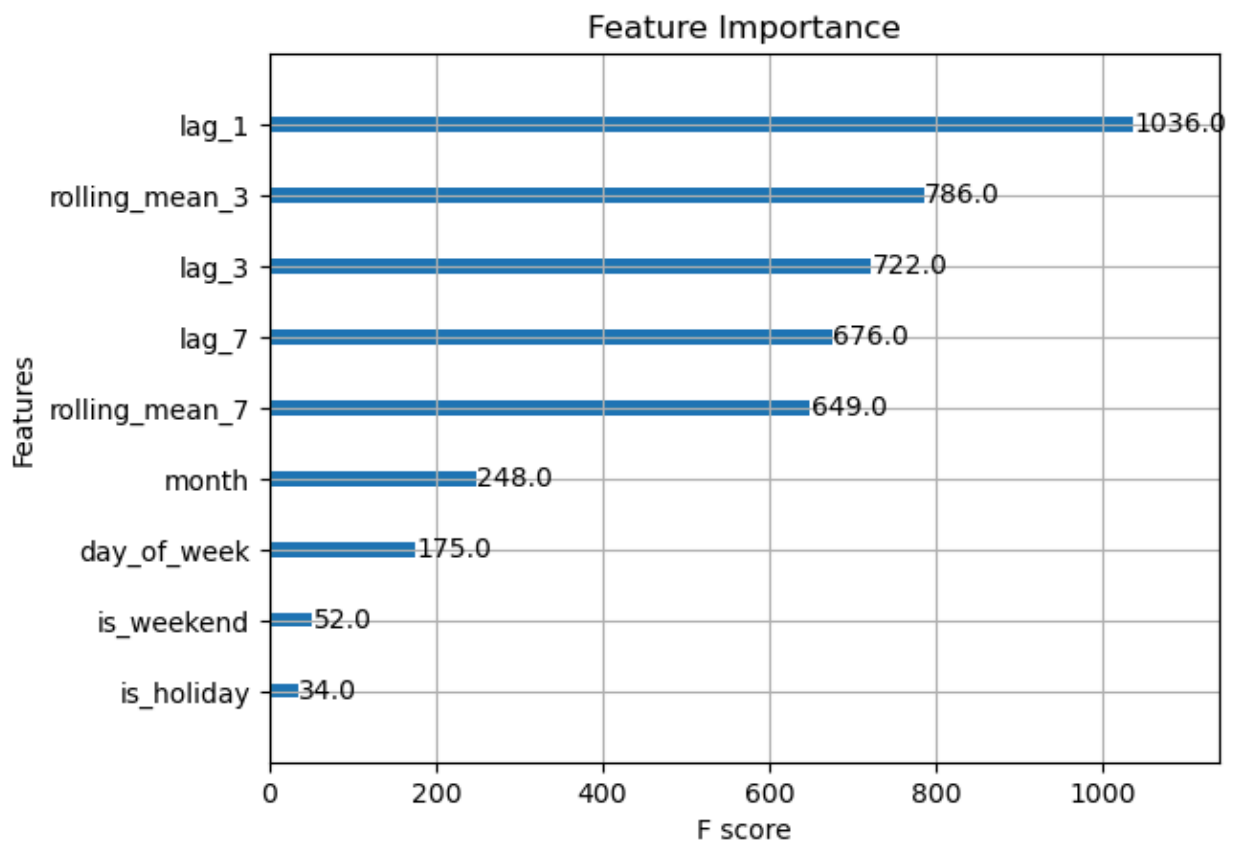
XGBRegressor
XGBRegressor(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None, early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None, feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=0.1, max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,

```

In [52]: *# Assuming xgb_model is our trained model*

```
plt.figure(figsize=(10, 6))
plot_importance(xgb_model)
plt.title("Feature Importance")
plt.show()
```

<Figure size 1000x600 with 0 Axes>



In [53]: *# Cross-Validation Using TimeSeriesSplit*

```
tscv = TimeSeriesSplit(n_splits=5)
model = XGBRegressor()

for train_index, test_index in tscv.split(X):
    X_train, X_test = X.iloc[train_index], X.iloc[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]

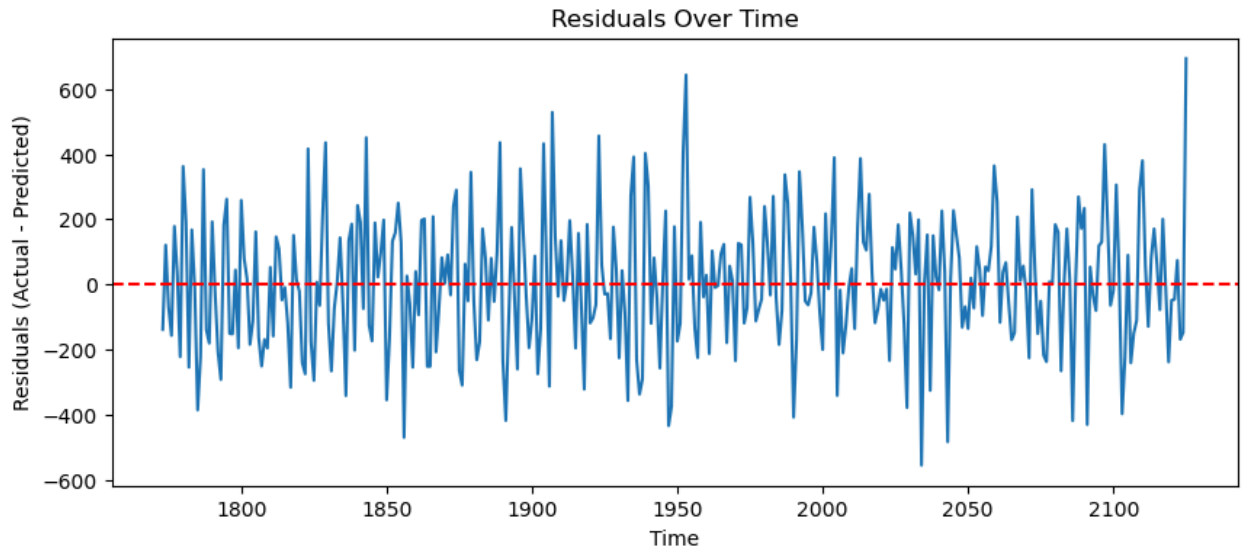
    model.fit(X_train, y_train)
    preds = model.predict(X_test)
    rmse = np.sqrt(mean_squared_error(y_test, preds))
    print(f"Fold RMSE: {rmse:.2f}")
```

Fold RMSE: 235.28
 Fold RMSE: 229.19
 Fold RMSE: 221.77
 Fold RMSE: 225.94
 Fold RMSE: 204.03

In [54]: *# Residual Plot (Error Analysis)*

```
residuals = y_test - preds
plt.figure(figsize=(10, 4))
plt.plot(residuals)
plt.title("Residuals Over Time")
plt.xlabel("Time")
```

```
plt.ylabel("Residuals (Actual - Predicted)")
plt.axhline(0, color='red', linestyle='--')
plt.show()
```



In [55]: *# Model Persistence*

```
# Save trained models for deployment
joblib.dump(xgb_model, "xgb_model.pkl")
```

Out[55]: ['xgb_model.pkl']

In [56]: *# Optional: add a version or timestamp*

```
timestamp = datetime.now().strftime("%Y%m%d_%H%M%S")
```

```
# Define filenames
```

```
xgb_filename = f"xgb_model_{timestamp}.pkl"
```

```
lr_filename = f"linear_model_{timestamp}.pkl" # Example if you trained another
```

```
# Save models
```

```
joblib.dump(xgb_model, xgb_filename)
```

```
# Optional: if you have other models
```

```
# joblib.dump(linear_model, lr_filename)
```

```
print(f"Models saved as: {xgb_filename}")
```

Models saved as: xgb_model_20250621_113318.pkl

In [57]: *# Evaluation Metrics Summary:*

```
# Calculate evaluation metrics
```

```
rmse = np.sqrt(mean_squared_error(y_test, preds))
```

```
mae = mean_absolute_error(y_test, preds)
```

```
r2 = r2_score(y_test, preds)
```

```
# Print the results
```

```
print(f"Root Mean Squared Error (RMSE): {rmse:.2f}")
print(f"Mean Absolute Error (MAE): {mae:.2f}")
print(f"R-squared (R²): {r2:.2f}")
```

Root Mean Squared Error (RMSE): 204.03

Mean Absolute Error (MAE): 162.57

R-squared (R²): 0.36

```
In [58]: # Save metrics to a file
metrics = {
    "RMSE": rmse,
    "MAE": mae,
    "R2": r2,
    "model_file": xgb_filename
}

import json
with open(f"metrics_{timestamp}.json", "w") as f:
    json.dump(metrics, f, indent=4)

print("Evaluation metrics saved.")
```

Evaluation metrics saved.

In []: